

In [4]:

*#Importing library*

```
import keras
from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout, Flatten, Conv2D, MaxPooling2D
from keras.layers.normalization import BatchNormalization
```

In [5]:

```
nn_random_seed(1000)
```

In [6]:

*#Instantiation*

```
AlexNet = Sequential()
```

*#1st Convolutional Layer*

```
AlexNet.add(Conv2D(filters=96, input_shape=(32,32,3), kernel_size=(11,11), strides=(4,4), padding='same'))
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='same'))
```

In [7]:

*#2nd Convolutional Layer*

```
AlexNet.add(Conv2D(filters=256, kernel_size=(5, 5), strides=(1,1), padding='same'))
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='same'))
```

In [8]:

*#3rd Convolutional Layer*

```
AlexNet.add(Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), padding='same'))
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))
```

In [9]:

*#4th Convolutional Layer*

```
AlexNet.add(Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), padding='same'))
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))
```

In [10]:

*#5th Convolutional Layer*

```
AlexNet.add(Conv2D(filters=256, kernel_size=(3,3), strides=(1,1), padding='same'))
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='same'))
```

In [11]:



```
#Passing it to a Fully Connected Layer
AlexNet.add(Flatten())
# 1st Fully Connected Layer
AlexNet.add(Dense(4096, input_shape=(32,32,3)))
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))
# Add Dropout to prevent overfitting
AlexNet.add(Dropout(0.4))
```

In [12]:



```
#2nd Fully Connected Layer
AlexNet.add(Dense(4096))
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))
#Add Dropout
AlexNet.add(Dropout(0.4))
```

In [13]:



```
#3rd Fully Connected Layer
AlexNet.add(Dense(1000))
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('relu'))
#Add Dropout
AlexNet.add(Dropout(0.4))
```

In [14]:



```
#Output Layer
AlexNet.add(Dense(10))
AlexNet.add(BatchNormalization())
AlexNet.add(Activation('softmax'))
```



In [15]:

```
#Model Summary
AlexNet.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 8, 8, 96)	34944
batch_normalization (Batch Normalization)	(None, 8, 8, 96)	384
activation (Activation)	(None, 8, 8, 96)	0
max_pooling2d (MaxPooling2D)	(None, 4, 4, 96)	0
conv2d_1 (Conv2D)	(None, 4, 4, 256)	614656
batch_normalization_1 (Batch Normalization)	(None, 4, 4, 256)	1024
activation_1 (Activation)	(None, 4, 4, 256)	0
max_pooling2d_1 (MaxPooling2D)	(None, 2, 2, 256)	0
conv2d_2 (Conv2D)	(None, 2, 2, 384)	885120
batch_normalization_2 (Batch Normalization)	(None, 2, 2, 384)	1536
activation_2 (Activation)	(None, 2, 2, 384)	0
conv2d_3 (Conv2D)	(None, 2, 2, 384)	1327488
batch_normalization_3 (Batch Normalization)	(None, 2, 2, 384)	1536
activation_3 (Activation)	(None, 2, 2, 384)	0
conv2d_4 (Conv2D)	(None, 2, 2, 256)	884992
batch_normalization_4 (Batch Normalization)	(None, 2, 2, 256)	1024
activation_4 (Activation)	(None, 2, 2, 256)	0
max_pooling2d_2 (MaxPooling2D)	(None, 1, 1, 256)	0
flatten (Flatten)	(None, 256)	0
dense (Dense)	(None, 4096)	1052672
batch_normalization_5 (Batch Normalization)	(None, 4096)	16384
activation_5 (Activation)	(None, 4096)	0
dropout (Dropout)	(None, 4096)	0
dense_1 (Dense)	(None, 4096)	16781312
batch_normalization_6 (Batch Normalization)	(None, 4096)	16384
activation_6 (Activation)	(None, 4096)	0

dropout_1 (Dropout)	(None, 4096)	0
dense_2 (Dense)	(None, 1000)	4097000
batch_normalization_7 (Batch Normalization)	(None, 1000)	4000
activation_7 (Activation)	(None, 1000)	0
dropout_2 (Dropout)	(None, 1000)	0
dense_3 (Dense)	(None, 10)	10010
batch_normalization_8 (Batch Normalization)	(None, 10)	40
activation_8 (Activation)	(None, 10)	0
=====		
Total params: 25,730,506		
Trainable params: 25,709,350		
Non-trainable params: 21,156		

In [16]:

# Compiling the model

AlexNet.compile(loss = keras.losses.categorical\_crossentropy, optimizer= 'adam', metrics=['

In [17]:

#Keras Library for CIFAR dataset

from keras.datasets import cifar10

(x\_train, y\_train),(x\_test, y\_test)=cifar10.load\_data()

#Train-validation-test split

from sklearn.model\_selection import train\_test\_split

x\_train,x\_val,y\_train,y\_val=train\_test\_split(x\_train,y\_train,test\_size=.3)

In [18]:

#Dimension of the CIFAR10 dataset

print((x\_train.shape,y\_train.shape))

print((x\_val.shape,y\_val.shape))

print((x\_test.shape,y\_test.shape))

((35000, 32, 32, 3), (35000, 1))

((15000, 32, 32, 3), (15000, 1))

((10000, 32, 32, 3), (10000, 1))

In [19]:

#Onehot Encoding the Labels.

from sklearn.utils.multiclass import unique\_labels

from keras.utils import to\_categorical

In [20]:

```
#Since we have 10 classes we should expect the shape[1] of y_train,y_val and y_test to chan  
y_train=to_categorical(y_train)  
y_val=to_categorical(y_val)  
y_test=to_categorical(y_test)
```

In [21]:

```
#Verifying the dimension after one hot encoding  
print((x_train.shape,y_train.shape))  
print((x_val.shape,y_val.shape))  
print((x_test.shape,y_test.shape))
```

```
((35000, 32, 32, 3), (35000, 10))  
((15000, 32, 32, 3), (15000, 10))  
((10000, 32, 32, 3), (10000, 10))
```

In [22]:

```
#Learning Rate Annealer  
from keras.callbacks import ReduceLROnPlateau  
lrr= ReduceLROnPlateau( monitor='val_acc', factor=.01, patience=3, min_lr=1e-5)
```

In [23]:

```
#Defining the parameters  
batch_size= 100  
epochs=100  
learn_rate=.001
```

In [35]:

```
#Training the model  
AlexNet.fit(x_train, y_train)
```

```
1094/1094 [=====] - 1021s 808ms/step - loss: 1.8078  
- accuracy: 0.3486
```

Out[35]:

```
<tensorflow.python.keras.callbacks.History at 0x1086a5784f0>
```



In [44]:

```

#Defining function for confusion matrix plot
def plot_confusion_matrix(y_true, y_pred, classes,
                           normalize=False,
                           title=None,
                           cmap=plt.cm.Blues):

    if not title:
        if normalize:
            title = 'Normalized confusion matrix'
        else:
            title = 'Confusion matrix, without normalization'

    # Compute confusion matrix
    cm = confusion_matrix(y_true, y_pred)
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

#Print Confusion matrix
    fig, ax = plt.subplots(figsize=(7,7))
    im = ax.imshow(cm, interpolation='nearest', cmap=cmap)
    ax.figure.colorbar(im, ax=ax)
    # We want to show all ticks...
    ax.set(xticks=np.arange(cm.shape[1]),
          yticks=np.arange(cm.shape[0]),
          xticklabels=classes, yticklabels=classes,
          title=title,
          ylabel='True label',
          xlabel='Predicted label')

    # Rotate the tick labels and set their alignment.
    plt.setp(ax.get_xticklabels(), rotation=45, ha="right",
              rotation_mode="anchor")
    # Loop over data dimensions and create text annotations.
    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i in range(cm.shape[0]):
        for j in range(cm.shape[1]):
            ax.text(j, i, format(cm[i, j], fmt),
                    ha="center", va="center",
                    color="white" if cm[i, j] > thresh else "black")
    fig.tight_layout()
    return ax

np.set_printoptions(precision=2)

#Making prediction
y_pred=AlexNet.predict_classes(x_test)
y_true=np.argmax(y_test,axis=1)

#Plotting the confusion matrix
from sklearn.metrics import confusion_matrix
confusion_mtx=confusion_matrix(y_true,y_pred)

class_names=['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship']


# Plotting non-normalized confusion matrix
plot_confusion_matrix(y_true, y_pred, classes = class_names,title = 'Confusion matrix, with

```

Confusion matrix, without normalization

Out[44]:

```
<AxesSubplot:title={'center':'Confusion matrix, without normalization'}, xlabel='Predicted label', ylabel='True label'>
```



In [45]:

# Task 2

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten
from keras.optimizers import Adam
from keras.layers.normalization import BatchNormalization
from keras.utils import np_utils
from keras.layers import Conv2D, MaxPooling2D, ZeroPadding2D, GlobalAveragePooling2D
from keras.layers.advanced_activations import LeakyReLU
from keras.preprocessing.image import ImageDataGenerator
```

In [47]:

```
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>)

11493376/11490434 [=====] - 1s 0us/step

In [48]:

```
X_train = X_train.reshape(X_train.shape[0], 28, 28, 1)
X_test = X_test.reshape(X_test.shape[0], 28, 28, 1)

X_train = X_train.astype('float32')
X_test = X_test.astype('float32')

X_train/=255
X_test/=255
```

In [49]:

```
number_of_classes = 10

Y_train = np_utils.to_categorical(y_train, number_of_classes)
Y_test = np_utils.to_categorical(y_test, number_of_classes)
```



In [50]:



```
model = Sequential()

model.add(Conv2D(32, (3, 3), input_shape=(28,28,1)))
model.add(BatchNormalization(axis=-1))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3)))
model.add(BatchNormalization(axis=-1))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2,2)))

model.add(Conv2D(64,(3, 3)))
model.add(BatchNormalization(axis=-1))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3)))
model.add(BatchNormalization(axis=-1))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2,2)))

model.add(Flatten())
# Fully connected layer
model.add(Dense(512))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(Dropout(0.2))
model.add(Dense(10))

model.add(Activation('softmax'))
```

In [51]:



```
model.compile(loss='categorical_crossentropy', optimizer=Adam(), metrics=['accuracy'])
```

In [52]:



```
gen = ImageDataGenerator(rotation_range=8, width_shift_range=0.08, shear_range=0.3,
                        height_shift_range=0.08, zoom_range=0.08)

test_gen = ImageDataGenerator()
```

In [53]:



```
train_generator = gen.flow(X_train, Y_train, batch_size=64)
test_generator = test_gen.flow(X_test, Y_test, batch_size=64)
```

In [54]:



```
model.fit_generator(train_generator, steps_per_epoch=60000//64, epochs=5,  
                    validation_data=test_generator, validation_steps=10000//64)
```

C:\Users\VRINDA\anaconda3\lib\site-packages\tensorflow\python\keras\engine\t raining.py:1844: UserWarning: `Model.fit\_generator` is deprecated and will b e removed in a future version. Please use `Model.fit`, which supports genera tors.

```
warnings.warn("`Model.fit_generator` is deprecated and '
```

Epoch 1/5

```
937/937 [=====] - 328s 328ms/step - loss: 0.2618 -  
accuracy: 0.9186 - val_loss: 0.0348 - val_accuracy: 0.9884
```

Epoch 2/5

```
937/937 [=====] - 279s 298ms/step - loss: 0.0586 -  
accuracy: 0.9815 - val_loss: 0.0342 - val_accuracy: 0.9881
```

Epoch 3/5

```
937/937 [=====] - 285s 304ms/step - loss: 0.0446 -  
accuracy: 0.9865 - val_loss: 0.0339 - val_accuracy: 0.9892
```

Epoch 4/5

```
937/937 [=====] - 307s 328ms/step - loss: 0.0392 -  
accuracy: 0.9885 - val_loss: 0.0377 - val_accuracy: 0.9874
```

Epoch 5/5

```
937/937 [=====] - 284s 303ms/step - loss: 0.0303 -  
accuracy: 0.9907 - val_loss: 0.0248 - val_accuracy: 0.9924
```

Out[54]:

```
<tensorflow.python.keras.callbacks.History at 0x108161f96d0>
```

In [ ]:

